

14.170: Programming for Economists

1/12/2009-1/16/2009

Melissa Dell

Matt Notowidigdo

Paul Schrimpf

Lecture 5, Large Data Sets in Stata + Numerical Precision

Overview

- This lecture is part wrap-up lecture, part “tips and tricks”
- Focus is on dealing with large data sets and on numerical precision
- Numerical precision
 - Introduction to binary representation
 - Equilibrating matrices
- Large data sets
 - How Stata represents data in memory
 - Speeding up code
 - Tips and tricks for large data sets

Numerical precision

- What the @&*%&\$!^ is going on here?

```
local a = 0.7 + 0.1  
local b = 0.8  
display (`a' == `b')
```

```
local a = 0.75 + 0.05  
local b = 0.8  
display (`a' == `b')
```

```
+ local a = 0.7 + 0.1  
+ local b = 0.8  
+ display (`a' == `b')  
0  
  
+ local a = 0.75 + 0.05  
+ local b = 0.8  
+ display (`a' == `b')  
1
```

Binary numbers

- Computers store numbers in base 2 (“bits”)

$$14_{10} = 1110_2$$
$$(14 = 2 + 4 + 8)$$

$$170_{10} = 10101010_2$$
$$(170 = 2 + 8 + 32 + 128)$$

How are decimals stored?

Binary numbers, con't

$$0.875_{10} = 0.111_2$$

$$(0.875 = 0.5 + 0.25 + 0.125)$$

$$0.80_{10} = 0.11001100\overline{1100}_2$$

$$0.70_{10} = 0.10110011\overline{0011}_2$$

$$0.10_{10} = 0.00011001\overline{1001}_2$$

$$0.75_{10} = 0.11_2$$

$$0.05_{10} = 0.00001100\overline{1100}_2$$

QUESTION: Is there a repeating decimal in base 10 that is not repeating in base 2?

Precision issues in Mata

```
mata
```

```
A = (1e10, 2e10 \ 2e-10, 3e-10)
```

```
A
```

```
rank(A)
```

```
luinv(A)
```

```
A_inv = (-3e-10, 2e10 \ 2e-10, -1e10)
```

```
I = A * A_inv
```

```
I
```

```
end
```

Precision issues in Mata

```
: A
      1      2
+-----+
1 | 1.00000e+10  2.00000e+10 |
2 | 2.00000e-10  3.00000e-10 |
+-----+

: rank(A)
1

: luinv(A)
[symmetric]
      1      2
+-----+
1 | .          |
2 | .          |
+-----+

: A_inv = (-3e-10, 2e10 \ 2e-10, -1e10)

: I = A * A_inv

: I
[symmetric]
      1      2
+-----+
1 | 1          |
2 | 0          1 |
+-----+
```


Precision issues in Mata

```
Mata
```

```
r = c = 0
```

```
A = (1e10, 2e10 \ 2e-10, 3e-10)
```

```
A
```

```
rank(A)
```

```
luinv(A, 1e-15)
```

```
_equilrc(A, r, c)
```

```
A
```

```
r
```

```
c
```

```
rank(A)
```

```
luinv(A)
```

```
c' : *luinv(A) : *r'
```

```
end
```

```

: luinv(A, 1e-15)
      1      2
+-----+
1 | -3.00000e-10  2.00000e+10 |
2 |  2.00000e-10 -1.00000e+10 |
+-----+

: _equilrc(A, r, c)

: A
[symmetric]
      1      2
+-----+
1 |   .75      |
2 |   1      1 |
+-----+

: r
      1
+-----+
1 | 5.00000e-11 |
2 | 3333333333  |
+-----+

: c
      1      2
+-----+
1 | 1.5      1 |
+-----+


: rank(A)
2

: c'*luinv(A):*r'
      1      2
+-----+
1 | -3.00000e-10  2.00000e+10 |
2 |  2.00000e-10 -1.00000e+10 |
+-----+

```

Precision issues in Mata

Large data sets in Stata



- Computer architecture overview
 - CPU: executes instructions
 - RAM (also called the “memory”): stores frequently-accessed data
 - DISK (“hard drive”): stores not-as-frequently used data
- RAM is accessed electronically; DISK is accessed mechanically (that’s why you can HEAR it). Thus DISK is several orders of magnitude slower than RAM.
- In Stata, if you *ever* have to access the disk, you’re pretty much dead. Stata was not written to deal with data sets that are larger than the available RAM. It expects the data set to fit in memory.
- So when you type “set memory XXXm”, make sure that you are not setting the value to be larger than the available RAM (some operating systems won’t even let you, anyway).
- For >20-30 GB of data, Stata is not recommended. Consider Matlab or SAS.

Large data sets in Stata, con't

- Don't keep re-creating the same variables over and over again
- “preserve” can really help or really hurt. Know when to use it and when to avoid it
- Don't estimate parameters you don't care about
- Lots of “if” and “in” commands could slow things down
- Create “1% sample” to develop and test code (to prevent unanticipated crashes after code has been running for hours)

```
clear
set seed 12345
set mem 2000m
set matsize 2000
set more off
set obs 5000
gen myn = _n
gen id = 1 + floor((_n - 1)/100)
sort id myn
by id: gen t = 1 + floor((_n - 1) / 5)
gen x = invnormal(uniform())
gen fe = invnormal(uniform())
sort id t myn
by id t: replace fe = fe[1]
gen y = 2 + x + fe + 100 * invnormal(uniform())

reg y x
xi i.id*i.t
reg y x _I*

summ t
gen idXt = id * (r(max) + 1) + t
areg y x, absorb(idXt)
```

Two-way fixed effects

```

. xi i.id*i.t
i.id      _Iid_1-50      (naturally coded; _Iid_1 omitted)
i.t       _It_1-20       (naturally coded; _It_1 omitted)
i.id*i.t  _IidXt_#_#     (coded as above)

```

```

. reg y x _I*

```

Source	SS	df	MS	Number of obs =	5000
Model	9217138.8	1000	9217.1388	F(1000, 3999) =	0.88
Residual	41898622.7	3999	10477.275	Prob > F =	0.9941
				R-squared =	0.1803
				Adj R-squared =	-0.0247
Total	51115761.5	4999	10225.1973	Root MSE =	102.36

y	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
x	.9800434	1.620583	0.60	0.545	-2.197203 4.15729
_Iid_2	15.14362	64.74499	0.23	0.815	-111.7926 142.0799
_Iid_3	49.41093	64.73921	0.76	0.445	-77.51401 176.3359
...
_Iid_49	44.54136	64.73895	0.69	0.491	-82.38307 171.4658
_Iid_50	25.32036	64.73848	0.39	0.696	-101.6031 152.2439
_It_2	-69.70522	64.74047	-1.08	0.282	-196.6326 57.22219
_It_3	-29.24825	64.73896	-0.45	0.651	-156.1727 97.67618
...
_It_19	-1.257793	64.74614	-0.02	0.985	-128.1963 125.6807
_It_20	28.69263	64.74172	0.44	0.658	-98.23723 155.6225
_IidXt_2_2	83.21349	91.55262	0.91	0.363	-96.28068 262.7077
_IidXt_2_3	.3550745	91.56491	0.00	0.997	-179.1632 179.8733
...
_IidXt_50_19	-33.42295	91.55229	-0.37	0.715	-212.9165 146.0706
_IidXt_50_20	-50.37966	91.55644	-0.55	0.582	-229.8813 129.122
_cons	-5.582456	45.77867	-0.12	0.903	-95.33416 84.16925

```

. areg y x, absorb(idXt)

```

Linear regression, absorbing indicators

```

Number of obs = 5000
F( 1, 3999) = 0.37
Prob > F = 0.5454
R-squared = 0.1803
Adj R-squared = -0.0247
Root MSE = 102.36

```

y	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
x	.9800434	1.620583	0.60	0.545	-2.197203 4.15729
_cons	2.549234	1.447597	1.76	0.078	-.2888639 5.387331
idXt	F(999, 3999) =		0.880	0.994	(1000 categories)

Two-way fixed effects

Fixed Effects with large data sets

```
clear
set seed 12345
set mem 100m
set more off
set obs 500000
gen myn = _n
gen id = 1 + floor((_n - 1)/200)
sort id myn
by id: gen t = _n
```

```
gen x = invnormal(uniform())
gen id_fe = invnormal(uniform())
gen t_fe = invnormal(uniform())
by id: replace id_fe = id_fe[1]
sort t id
by t: replace t_fe = t_fe[1]
gen y = 2 + x + id_fe + t_fe + 100 * invnormal(uniform())

xi i.t
xtreg y x _It*, i(id) fe
```

~674 seconds

Fixed Effects with large data sets

```
. xi i,t
i,t          _It_1-200          (naturally coded; _It_1 omitted)
```

```
. xtreg y x _It*, i(id) fe
```

```
Fixed-effects (within) regression          Number of obs   =   500000
Group variable: id                        Number of groups  =    2500

R-sq:  within = 0.0008                    Obs per group:   min =    200
      between = 0.0009                      avg   =   200.0
      overall  = 0.0008                      max   =    200

                                           F(200,497300)    =    1.87
corr(u_i, Xb)  = 0.0006                    Prob > F         =    0.0000
```

	y	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
	x	1.224538	.1416581	8.64	0.000	.9468925	1.502184
	_It_2	.5291215	2.827906	0.19	0.852	-5.013486	6.071729
	_It_3	.7475153	2.827904	0.26	0.792	-4.795089	6.29012
	_It_4	2.120499	2.827907	0.75	0.453	-3.42211	7.663107
	_It_5	.1249969	2.827904	0.04	0.965	-5.417607	5.667601
	_It_6	-.5349088	2.827912	-0.19	0.850	-6.077528	5.00771
	_It_7	-.90349	2.827906	-0.32	0.749	-6.446097	4.639117
	_It_8	-.9549379	2.827904	-0.34	0.736	-6.497512	4.587666

Fixed Effects with large data sets

```
clear
set seed 12345
set mem 100m
set more off
set obs 500000
gen myn = _n
gen id = 1 + floor((_n - 1)/200)
sort id myn
by id: gen t = _n
```

```
gen x = invnormal(uniform())
gen id_fe = invnormal(uniform())
gen t_fe = invnormal(uniform())
by id: replace id_fe = id_fe[1]
sort t id
by t: replace t_fe = t_fe[1]
gen y = 2 + x + id_fe + t_fe + 100 * invnormal(uniform())
```

```
xtreg y, i(id) fe
predict y_resid, e
xtreg x, i(id) fe
predict x_resid, e
xtreg y_resid x_resid, i(t) fe
```

~53 seconds

Fixed Effects with large data sets

```
. xtreg y_resid x_resid, i(t) fe
warning: existing panel variable is not t
```

Fixed-effects (within) regression
Group variable: t

Number of obs = 500000
Number of groups = 200

R-sq: within = 0.0002
between = 0.0030
overall = 0.0002

Obs per group: min = 2500
avg = 2500.0
max = 2500

corr(u_i, Xb) = 0.0008

F(1,499799) = 75.10
Prob > F = 0.0000

y_resid	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
x_resid	1.224538	.1413035	8.67	0.000	.9475875	1.501489
_cons	-2.89e-15	.1410413	-0.00	1.000	-.2764365	.2764365
sigma_u	2.4508088					
sigma_e	99.731239					
rho	.00060352	(fraction of variance due to u_i)				
F test that all u_i=0:		F(199, 499799) =	1.51	Prob > F = 0.0000		

Other tips and tricks when you have large number of fixed effects in large data sets

- Use matrix algebra
- Newton steps in parallel
- “zig-zag maximization”
(Heckman-McCurdy)

Matrix algebra

```
clear mata
```

```
mata
```

```
rseed(14170)
```

```
N = 3000
```

```
rA = rnormal(5, 5, 0, 1)
```

```
rB = rnormal(5, N, 0, 1)
```

```
rC = rnormal(N, 5, 0, 1)
```

```
d = rnormal(1, N, 0, 1)
```

```
V = (rA, rB \ rC, diag(d))
```

```
V_inv = luinv(V)
```

```
V_inv[1..5, 1..5]
```

~162 seconds

Matrix algebra

```
clear mata
mata
rseed(14170)
N = 3000
rA = rnormal(5, 5, 0, 1)
rB = rnormal(5, N, 0, 1)
rC = rnormal(N, 5, 0, 1)
d = rnormal(1, N, 0, 1)
V = (rA, rB \ rC, diag(d))
V_fast = luinv(rA - cross(rB', d :^ -1, rC))
V_fast
```

<1 second

$$\begin{bmatrix} \mathbf{A} & \mathbf{B} \\ \mathbf{C} & \mathbf{D} \end{bmatrix}^{-1} = \begin{bmatrix} (\mathbf{A} - \mathbf{B}\mathbf{D}^{-1}\mathbf{C})^{-1} & -(\mathbf{A} - \mathbf{B}\mathbf{D}^{-1}\mathbf{C})^{-1}\mathbf{B}\mathbf{D}^{-1} \\ -\mathbf{D}^{-1}\mathbf{C}(\mathbf{A} - \mathbf{B}\mathbf{D}^{-1}\mathbf{C})^{-1} & \mathbf{D}^{-1} + \mathbf{D}^{-1}\mathbf{C}(\mathbf{A} - \mathbf{B}\mathbf{D}^{-1}\mathbf{C})^{-1}\mathbf{B}\mathbf{D}^{-1} \end{bmatrix}$$

Fixed Effects probit

- Finkelstein, Luttmer, Notowidigdo (2008) run Fixed Effects probit as a robustness check
 - What about the incidental parameters problem? (see Hahn and Newey, EMA, 2004)
- But what to do with $>11,000$ fixed effects!
 - Cannot de-mean within panel as you could with linear probability model
 - Stata/SE and Stata/MP matrix size limit is 11,000
 - Need several computation tricks

Fixed Effects probit

```
clear
set seed 12345
set matsize 2000
set obs 2000
gen id = 1+floor((_n - 1)/4)
gen a = invnormal(uniform())
gen fe_raw = 0.5*invnorm(uniform()) + 2*a
bys id: egen fe = mean(fe_raw)
gen x = invnormal(uniform())
gen e = invnormal(uniform())
gen y = (1*x + fe > invnormal(uniform())) + a)

bys id: egen x_mean = mean(x)
gen x_demean = x - x_mean
probit y x
probit y x_demean
sort id y
by id: keep if y[1] != y[_N]
probit y x
xi i.id
probit y x _I*
```

Fixed Effects probit

```
. probit y x
```

```
Iteration 0: log likelihood = -1386.2304
Iteration 1: log likelihood = -1175.0473
Iteration 2: log likelihood = -1169.7587
Iteration 3: log likelihood = -1169.7486
```

```
Probit regression               Number of obs   =       2000
                               LR chi2(1)        =       432.96
                               Prob > chi2        =       0.0000
Log likelihood = -1169.7486     Pseudo R2      =       0.1562
```

	y	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]
x		.6681063	.0350562	19.06	0.000	.5993974 .7368153
_cons		.01189	.0300833	0.40	0.693	-.0470722 .0708522

```
. probit y x_demean
```

```
Iteration 0: log likelihood = -1386.2304
Iteration 1: log likelihood = -1222.2689
Iteration 2: log likelihood = -1219.6961
Iteration 3: log likelihood = -1219.6943
```

```
Probit regression               Number of obs   =       2000
                               LR chi2(1)        =       333.07
                               Prob > chi2        =       0.0000
Log likelihood = -1219.6943     Pseudo R2      =       0.1201
```

	y	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]
x_demean		.6597479	.0385849	17.10	0.000	.5841229 .7353729
_cons		.0136155	.0295461	0.46	0.645	-.0442938 .0715249

```
. probit y x
```

```
Iteration 0: log likelihood = -1181.104
Iteration 1: log likelihood = -996.25035
Iteration 2: log likelihood = -991.28997
Iteration 3: log likelihood = -991.27891
```

```
Probit regression               Number of obs   =       1704
                               LR chi2(1)        =       379.65
                               Prob > chi2        =       0.0000
Log likelihood = -991.27891     Pseudo R2      =       0.1607
```

	y	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]
x		.6765638	.0380693	17.77	0.000	.6019494 .7511782
_cons		.0089522	.0326835	0.27	0.784	-.0551062 .0730106

```
. xi i,id
i,id          _Iid_1-500      (naturally coded; _Iid_1 omitted)
```

```
. probit y x _I*
```

```
Iteration 0: log likelihood = -1181.104
Iteration 1: log likelihood = -838.75262
Iteration 2: log likelihood = -805.34791
Iteration 3: log likelihood = -803.269
Iteration 4: log likelihood = -803.25405
Iteration 5: log likelihood = -803.25404
```

```
Probit regression               Number of obs   =       1704
                               LR chi2(426)      =       755.70
                               Prob > chi2        =       0.0000
Log likelihood = -803.25404     Pseudo R2      =       0.3199
```

	y	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]
x		1.024471	.0564683	18.14	0.000	.9137949 1.135146
_Iid_2		-.5300751	1.028931	-0.52	0.606	-2.546743 1.486593
_Iid_3		1.293332	1.018006	1.27	0.204	-.7019239 3.288587
_Iid_4		-1.504805	1.184503	-1.27	0.204	-3.826388 .8167775
...
_Iid_499		-.5475299	1.0869	-0.50	0.614	-2.677815 1.582755
_Iid_500		-.1185591	1.10614	-0.11	0.915	-2.286554 2.049435
_cons		-.1931788	.7153942	-0.27	0.787	-1.595326 1.208968

Fixed Effects probit (slow)

```
clear
set more off
set mem 1000m
set seed 12345
set matsize 3000
set obs 12000
gen id = 1+floor((_n - 1)/4)
gen a = invnormal(uniform())
gen fe_raw = 0.5*invnorm(uniform()) + 2*a
bys id: egen fe = mean(fe_raw)
gen x = invnormal(uniform())
gen e = invnormal(uniform())
gen y = (1*x + fe > invnormal(uniform()) + a)

sort id y
by id: keep if y[1] != y[_N]

xi i.id
probit y x _I*
```

Fixed Effects probit (slow)

~40 minutes

```
*
. xi i.id
i.id          _Iid_1-3000      (naturally coded; _Iid_1 omitted)
```

```
. probit y x _I*
```

```
Iteration 0:  log likelihood = -7131.0824
Iteration 1:  log likelihood = -5185.5403
Iteration 2:  log likelihood = -5019.9566
Iteration 3:  log likelihood = -5011.8605
Iteration 4:  log likelihood = -5011.8245
Iteration 5:  log likelihood = -5011.8245
```

Probit regression

```
Number of obs   =      10288
LR chi2(2572)   =      4238.52
Prob > chi2      =      0.0000
Pseudo R2       =      0.2972
```

Log likelihood = -5011.8245

y	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
x	.9416602	.0213752	44.05	0.000	.8997655	.9835548
_Iid_2	.3597163	.9682755	0.37	0.710	-1.538069	2.257501
_Iid_4	-.0921325	.9684319	-0.10	0.924	-1.990224	1.805959
...
_Iid_2998	1.128068	.9542554	1.18	0.237	-.742238	2.998374
_Iid_2999	.5476889	.9590439	0.57	0.568	-1.332003	2.42738
_Iid_3000	.4456436	.918257	0.49	0.627	-1.354107	2.245394
_cons	-.6380591	.6898159	-0.92	0.355	-1.990073	.7139553

```

clear
set mem 1000m
set seed 12345
set matsize 3000
set obs 12000
gen id = 1+floor((_n - 1)/4)
gen a = invnormal(uniform())
gen fe_raw = 0.5*invnorm(uniform()) + 2*a
bys id: egen fe = mean(fe_raw)
gen x = invnormal(uniform())
gen e = invnormal(uniform())
gen y = (1*x + fe > invnormal(uniform())) + a)
sort id y
by id: keep if y[1] != y[_N]

egen id_new = group(id)
summ id_new
local max = r(max)
gen fe_hat = 0
forvalues iter = 1/20 {
    probit y x, nocons offset(fe_hat)
    capture drop xb*
    predict xb, xb nooffset
    forvalues i = 1/`max' {
        qui probit y if id_new == `i', offset(xb)
        qui replace fe_hat = _b[_cons] if id_new == `i'
    }
}
probit y x, noconstant offset(fe_hat)

```

Fixed Effects probit (faster)

Fixed Effects probit (faster)

~8 minutes

```
. probit y x, nocons offset(fe_hat)
```

```
Iteration 0:  log likelihood = -6936.1812  
Iteration 1:  log likelihood = -5089.0776  
Iteration 2:  log likelihood = -5012.1796  
Iteration 3:  log likelihood = -5011.8245
```

Probit regression

Log likelihood = -5011.8245

```
Number of obs   =    10288  
Wald chi2(1)    =    3073.53  
Prob > chi2     =     0.0000
```

y		Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
x		.9416602	.0169854	55.44	0.000	.9083694	.9749509
fe_hat		(offset)					

QUESTION: Why are standard errors not the same?

Exercises

- (A) Speed up fixed effects probit even more by updating fixed effects in parallel
- (B) Fix standard errors in FE probit example